

# Digging Deeper with PostgreSQL

**Breaking up processes and  
understanding performance**

# Agenda

- Case Study (10m)
- Key Concepts Review (10m)
- Advanced Functions (20m)
- Database Setup (5m)
- Explain and performance review with PGAdmin 4 (20m)
- Organizing: With, sub-select, temporary tables and views (1.5h)
  - Performance reviews included
- Indexing (30m)

# Case Study Overview

- Shopping experience
- Primary keys will always be 'id' column in tables
- All foreign keys will refer to: <table\_name>\_id

# Case Study: Shopping site

## **Shopping Components**

Products  
Product Types  
Users  
Orders  
Carts

# Case Study: Shopping site

## Shopping Components

- Products
- Product Types
- Users
- Orders
- Carts

### products

- id
- product\_type\_id
- price
- name
- grams
- address

# Case Study: Shopping site

## Shopping Components

Products

Product Types

Users

Orders

Carts

### **products**

id  
product\_type\_id  
price  
name  
grams  
address

### **product\_types**

id  
name  
description

# Case Study: Shopping site

## Shopping Components

- Products
- Product Types
- Users
- Orders
- Carts

### **products**

- id
- product\_type\_id
- price
- name
- grams
- address

### **product\_types**

- id
- name
- description



# Case Study: Shopping site

## Shopping Components

- Products
- Product Types
- Users
- Orders
- Carts

### products

- id
- product\_type\_id
- price
- name
- grams
- address

### product\_types

- id
- name
- description

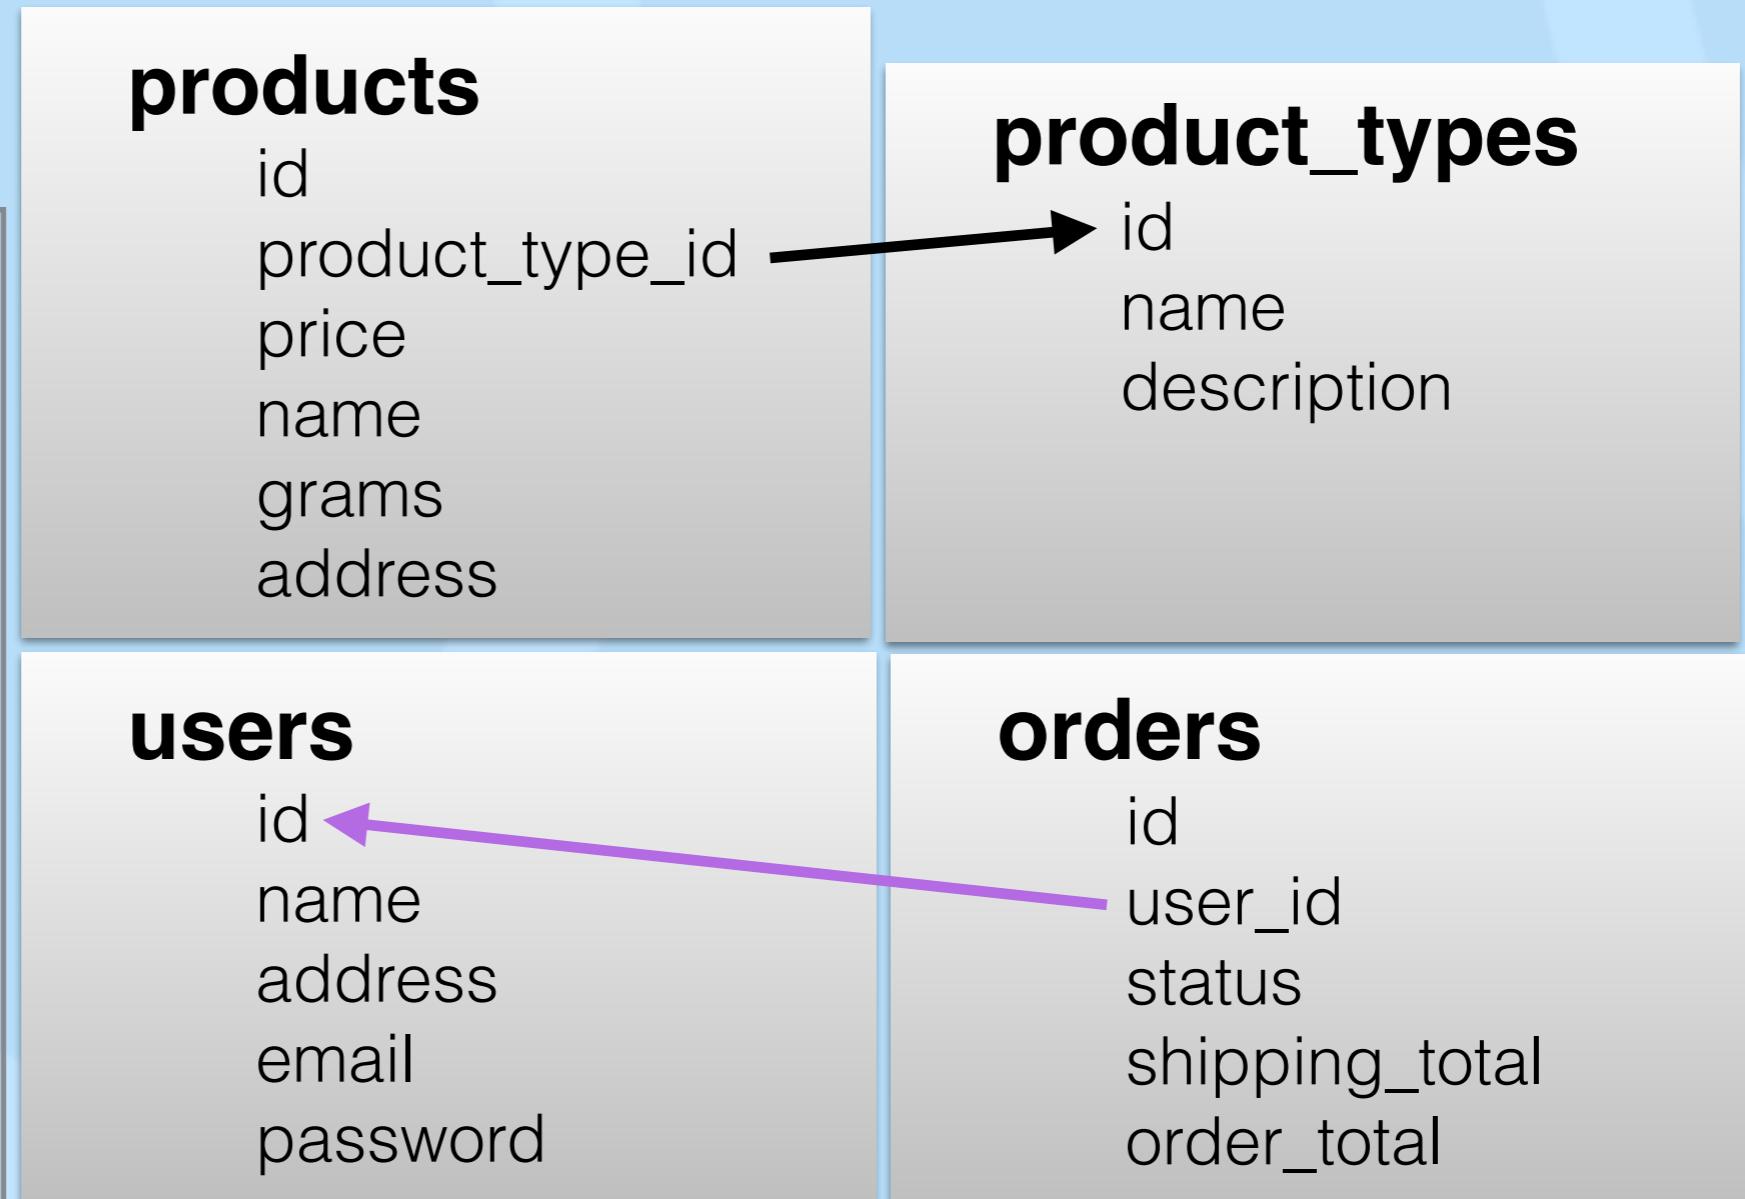
### users

- id
- name
- address
- email
- password

# Case Study: Shopping site

## Shopping Components

- Products
- Product Types
- Users
- Orders
- Carts



# Case Study: Shopping site

## Shopping Components

- Products
- Product Types
- Users
- Orders
- Carts



# Case Study: Shopping site

## Shopping Components

- Products
- Product Types
- Users
- Orders
- Carts



# Case Study: Shopping site

## Shopping Components

- Products
- Product Types
- Users
- Orders
- Carts



# Case Study: Shopping site

## products

id  
product\_type\_id  
price  
name  
grams  
address

## product\_types

id  
name  
description

## users

id  
name  
address  
email  
password

## orders

id  
user\_id  
status  
shipping\_total  
order\_total

## carts

id  
user\_id  
last\_updated

# Case Study: Shopping site

## products

id  
product\_type\_id  
price  
name  
grams  
address

## product\_types

id  
name  
description

## users

id  
name  
address  
email  
password

## orders

id  
user\_id  
status  
shipping\_total  
order\_total

## order\_products

id  
product\_id  
order\_id

## carts

id  
user\_id  
last\_updated

# Case Study: Shopping site

## products

id ←  
product\_type\_id →  
price  
name  
grams  
address

## product\_types

id  
name  
description

## users

id ←  
name  
address  
email  
password

## orders

id ←  
user\_id  
status  
shipping\_total  
order\_total

## order\_products

id  
product\_id  
order\_id

## carts

id  
user\_id  
last\_updated

# Case Study: Shopping site

## products

id ←  
product\_type\_id →  
price  
name  
grams  
address

## product\_types

id  
name  
description

## users

id ←  
name  
address  
email  
password

## orders

id ←  
user\_id  
status  
shipping\_total  
order\_total

## order\_products

id  
product\_id  
order\_id

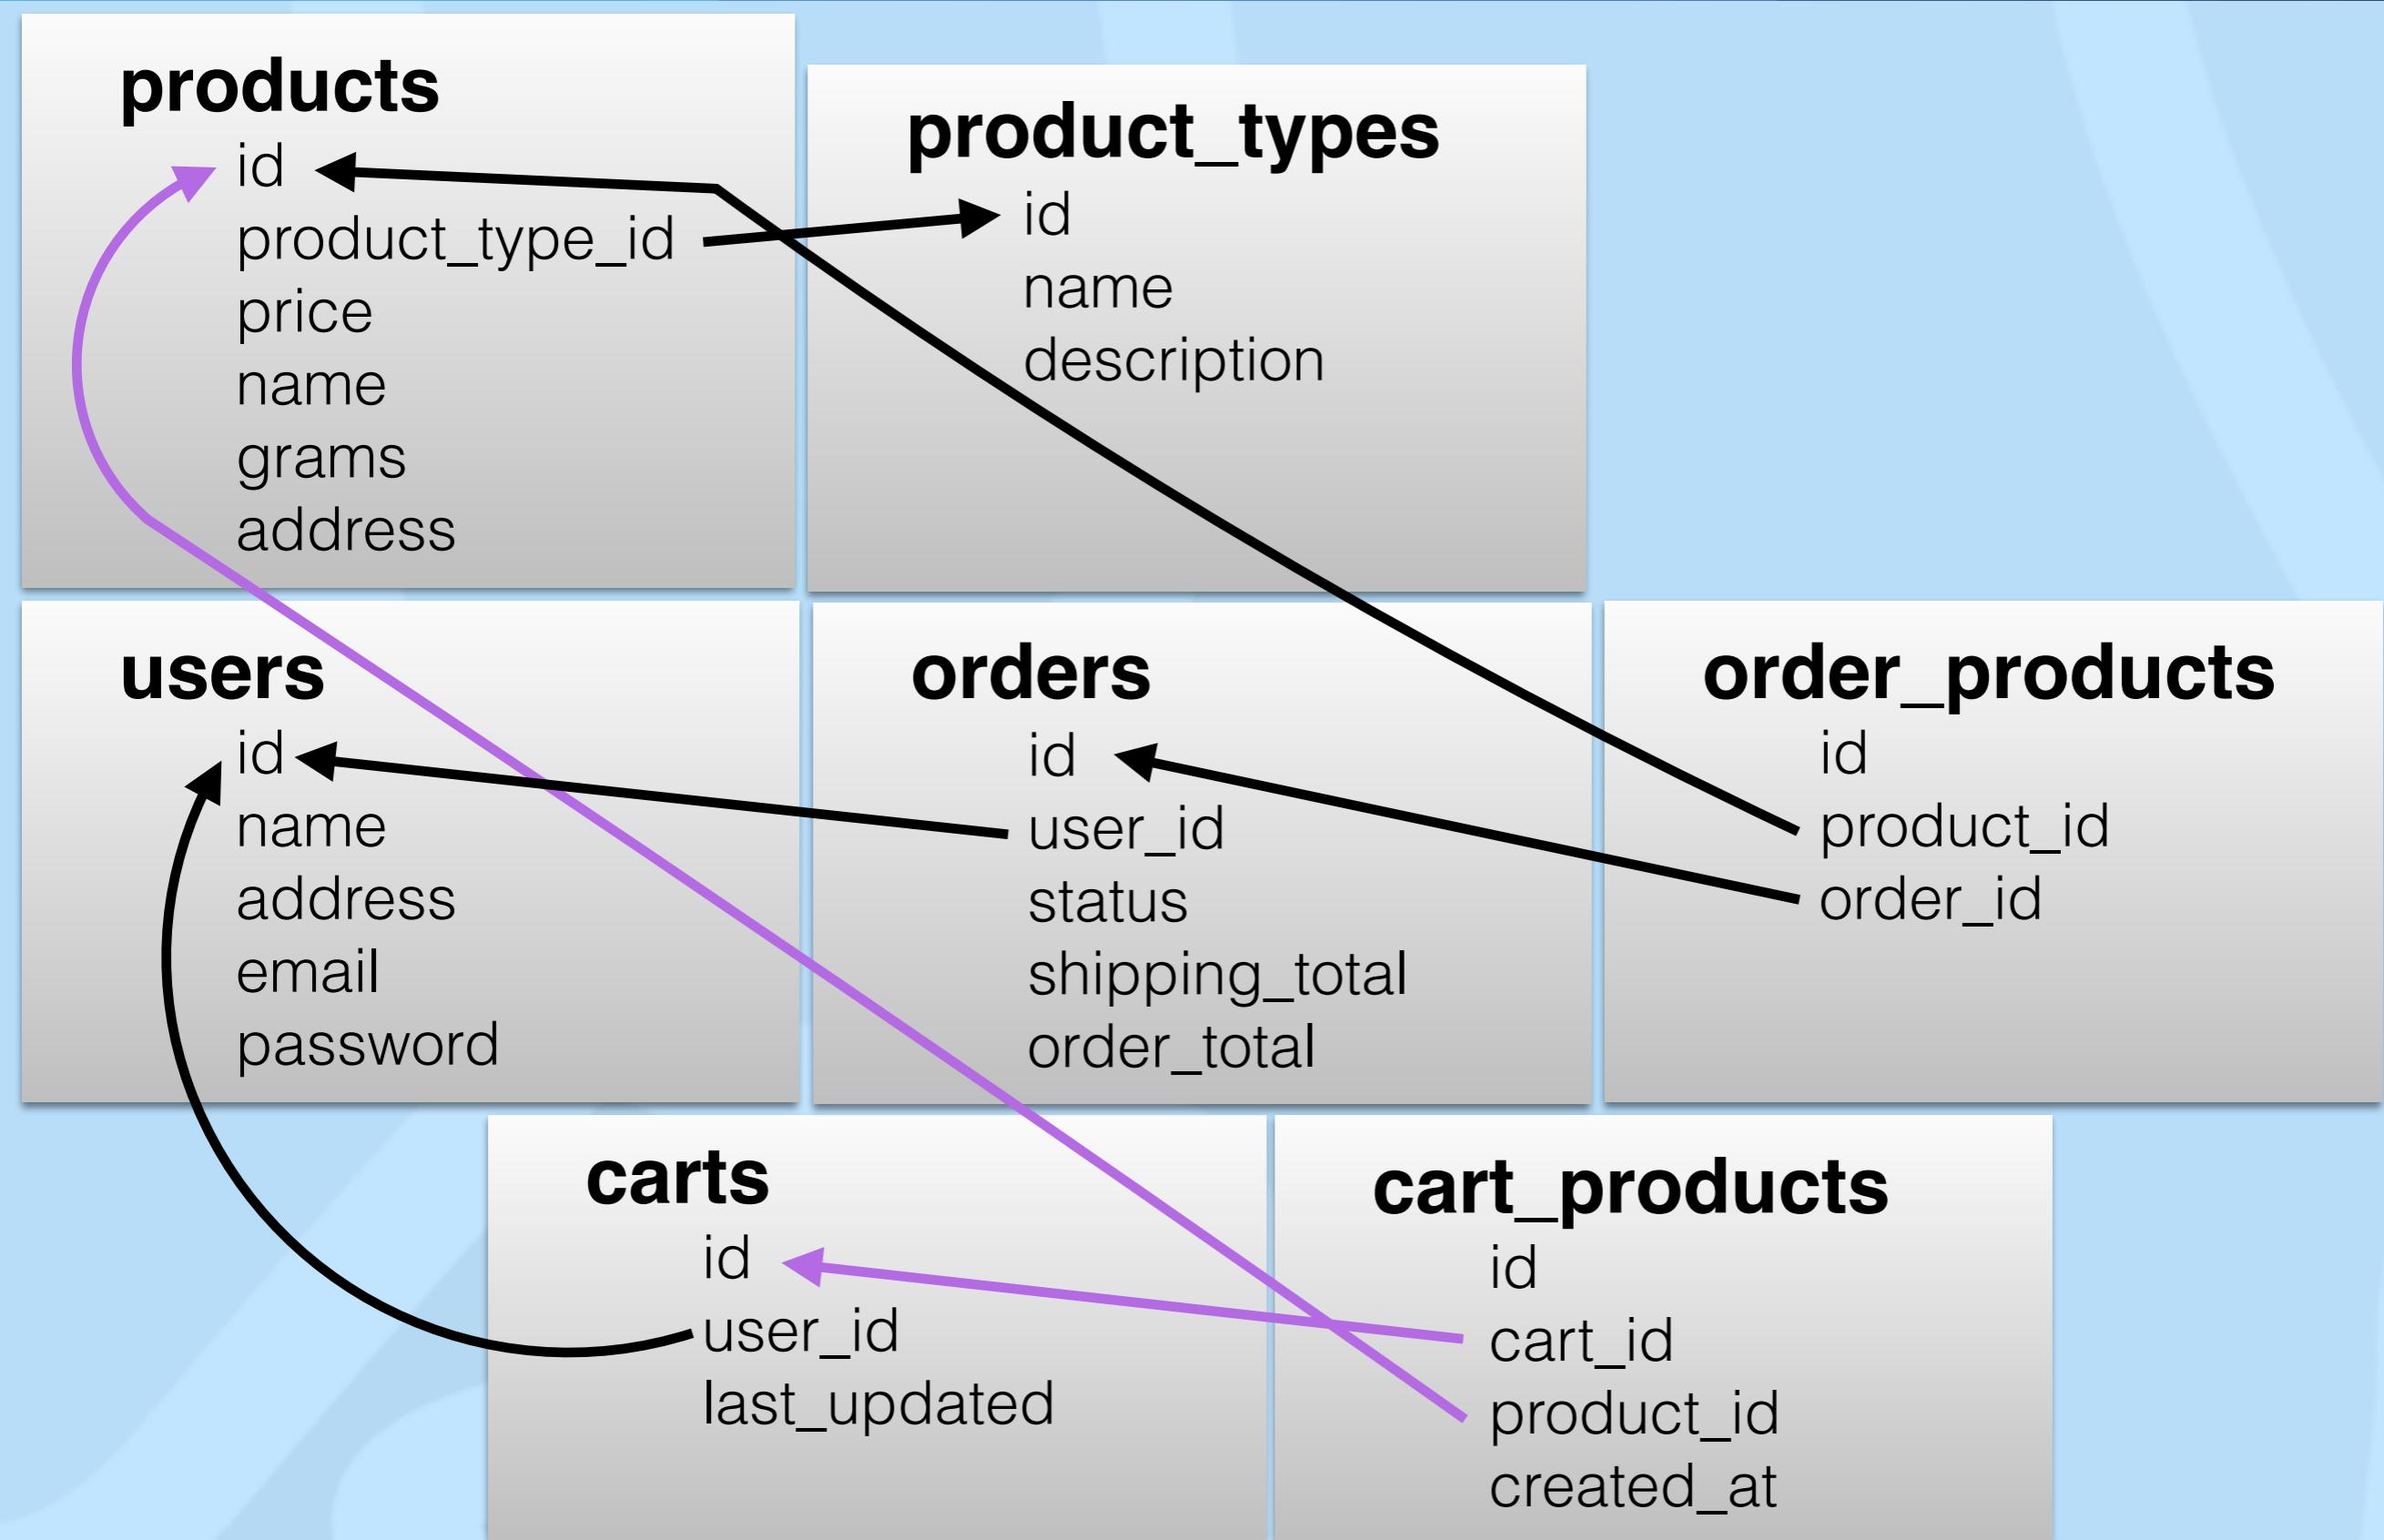
## carts

id  
user\_id  
last\_updated

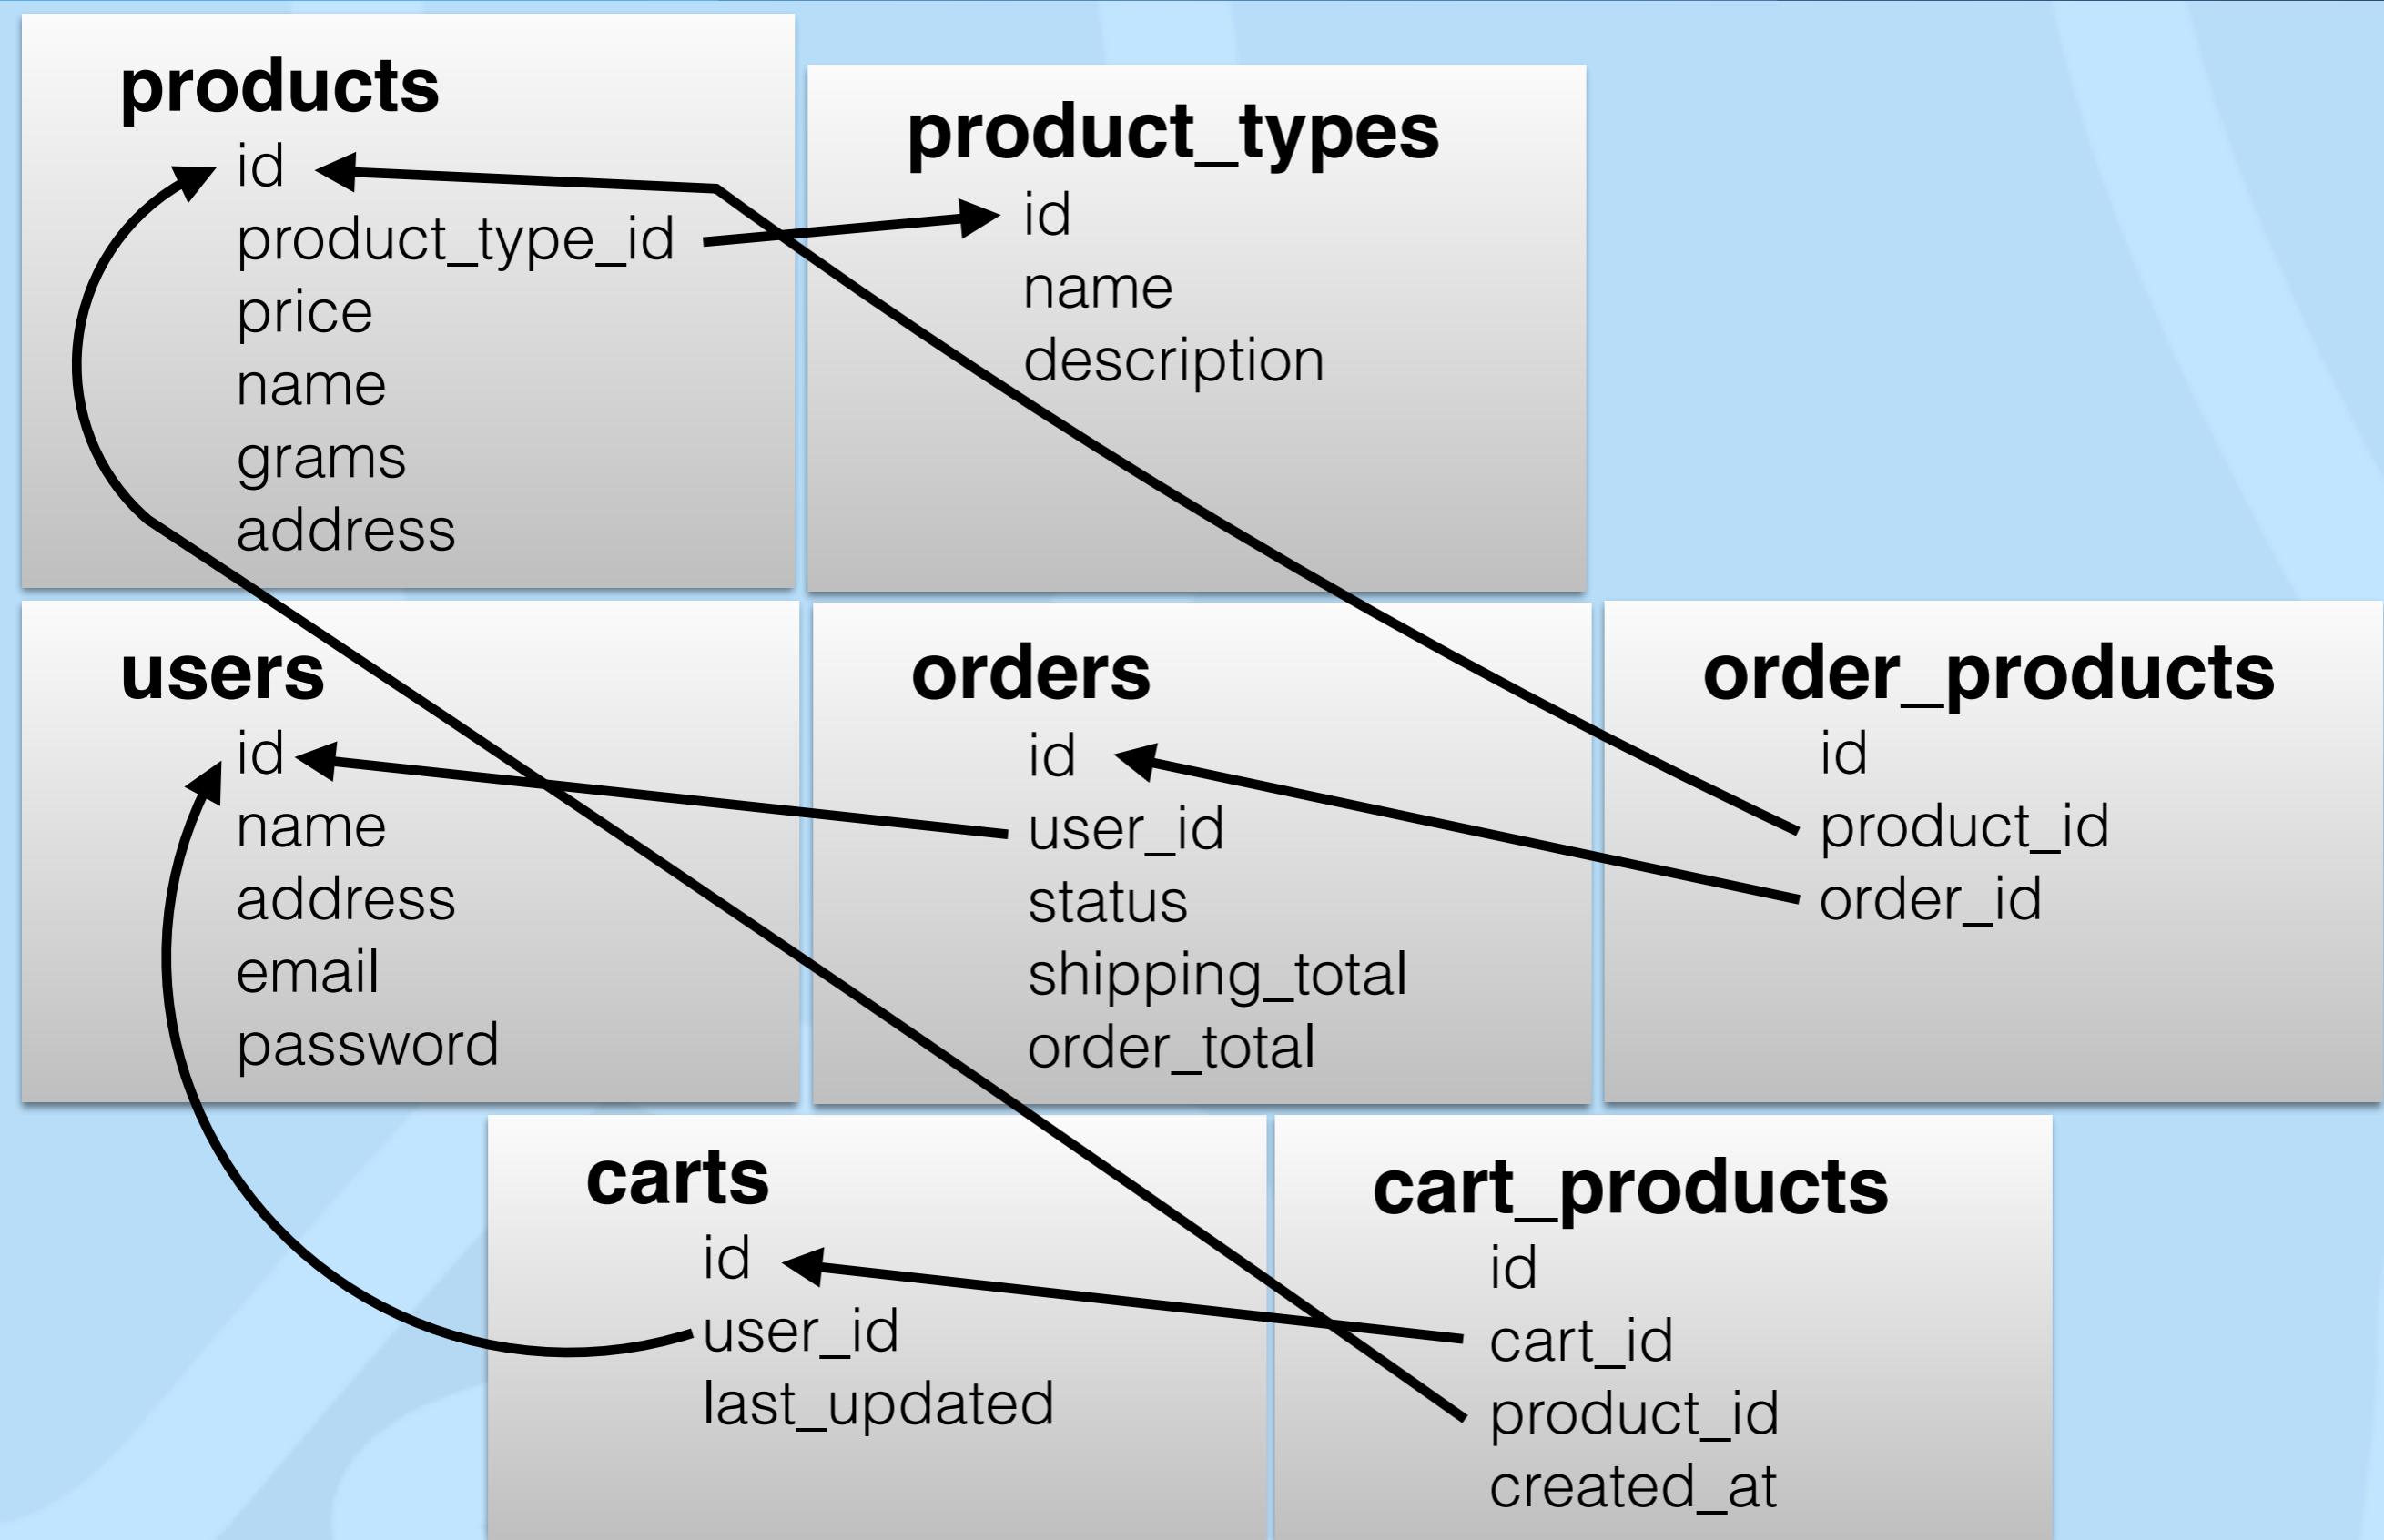
## cart\_products

id  
cart\_id  
product\_id  
created\_at

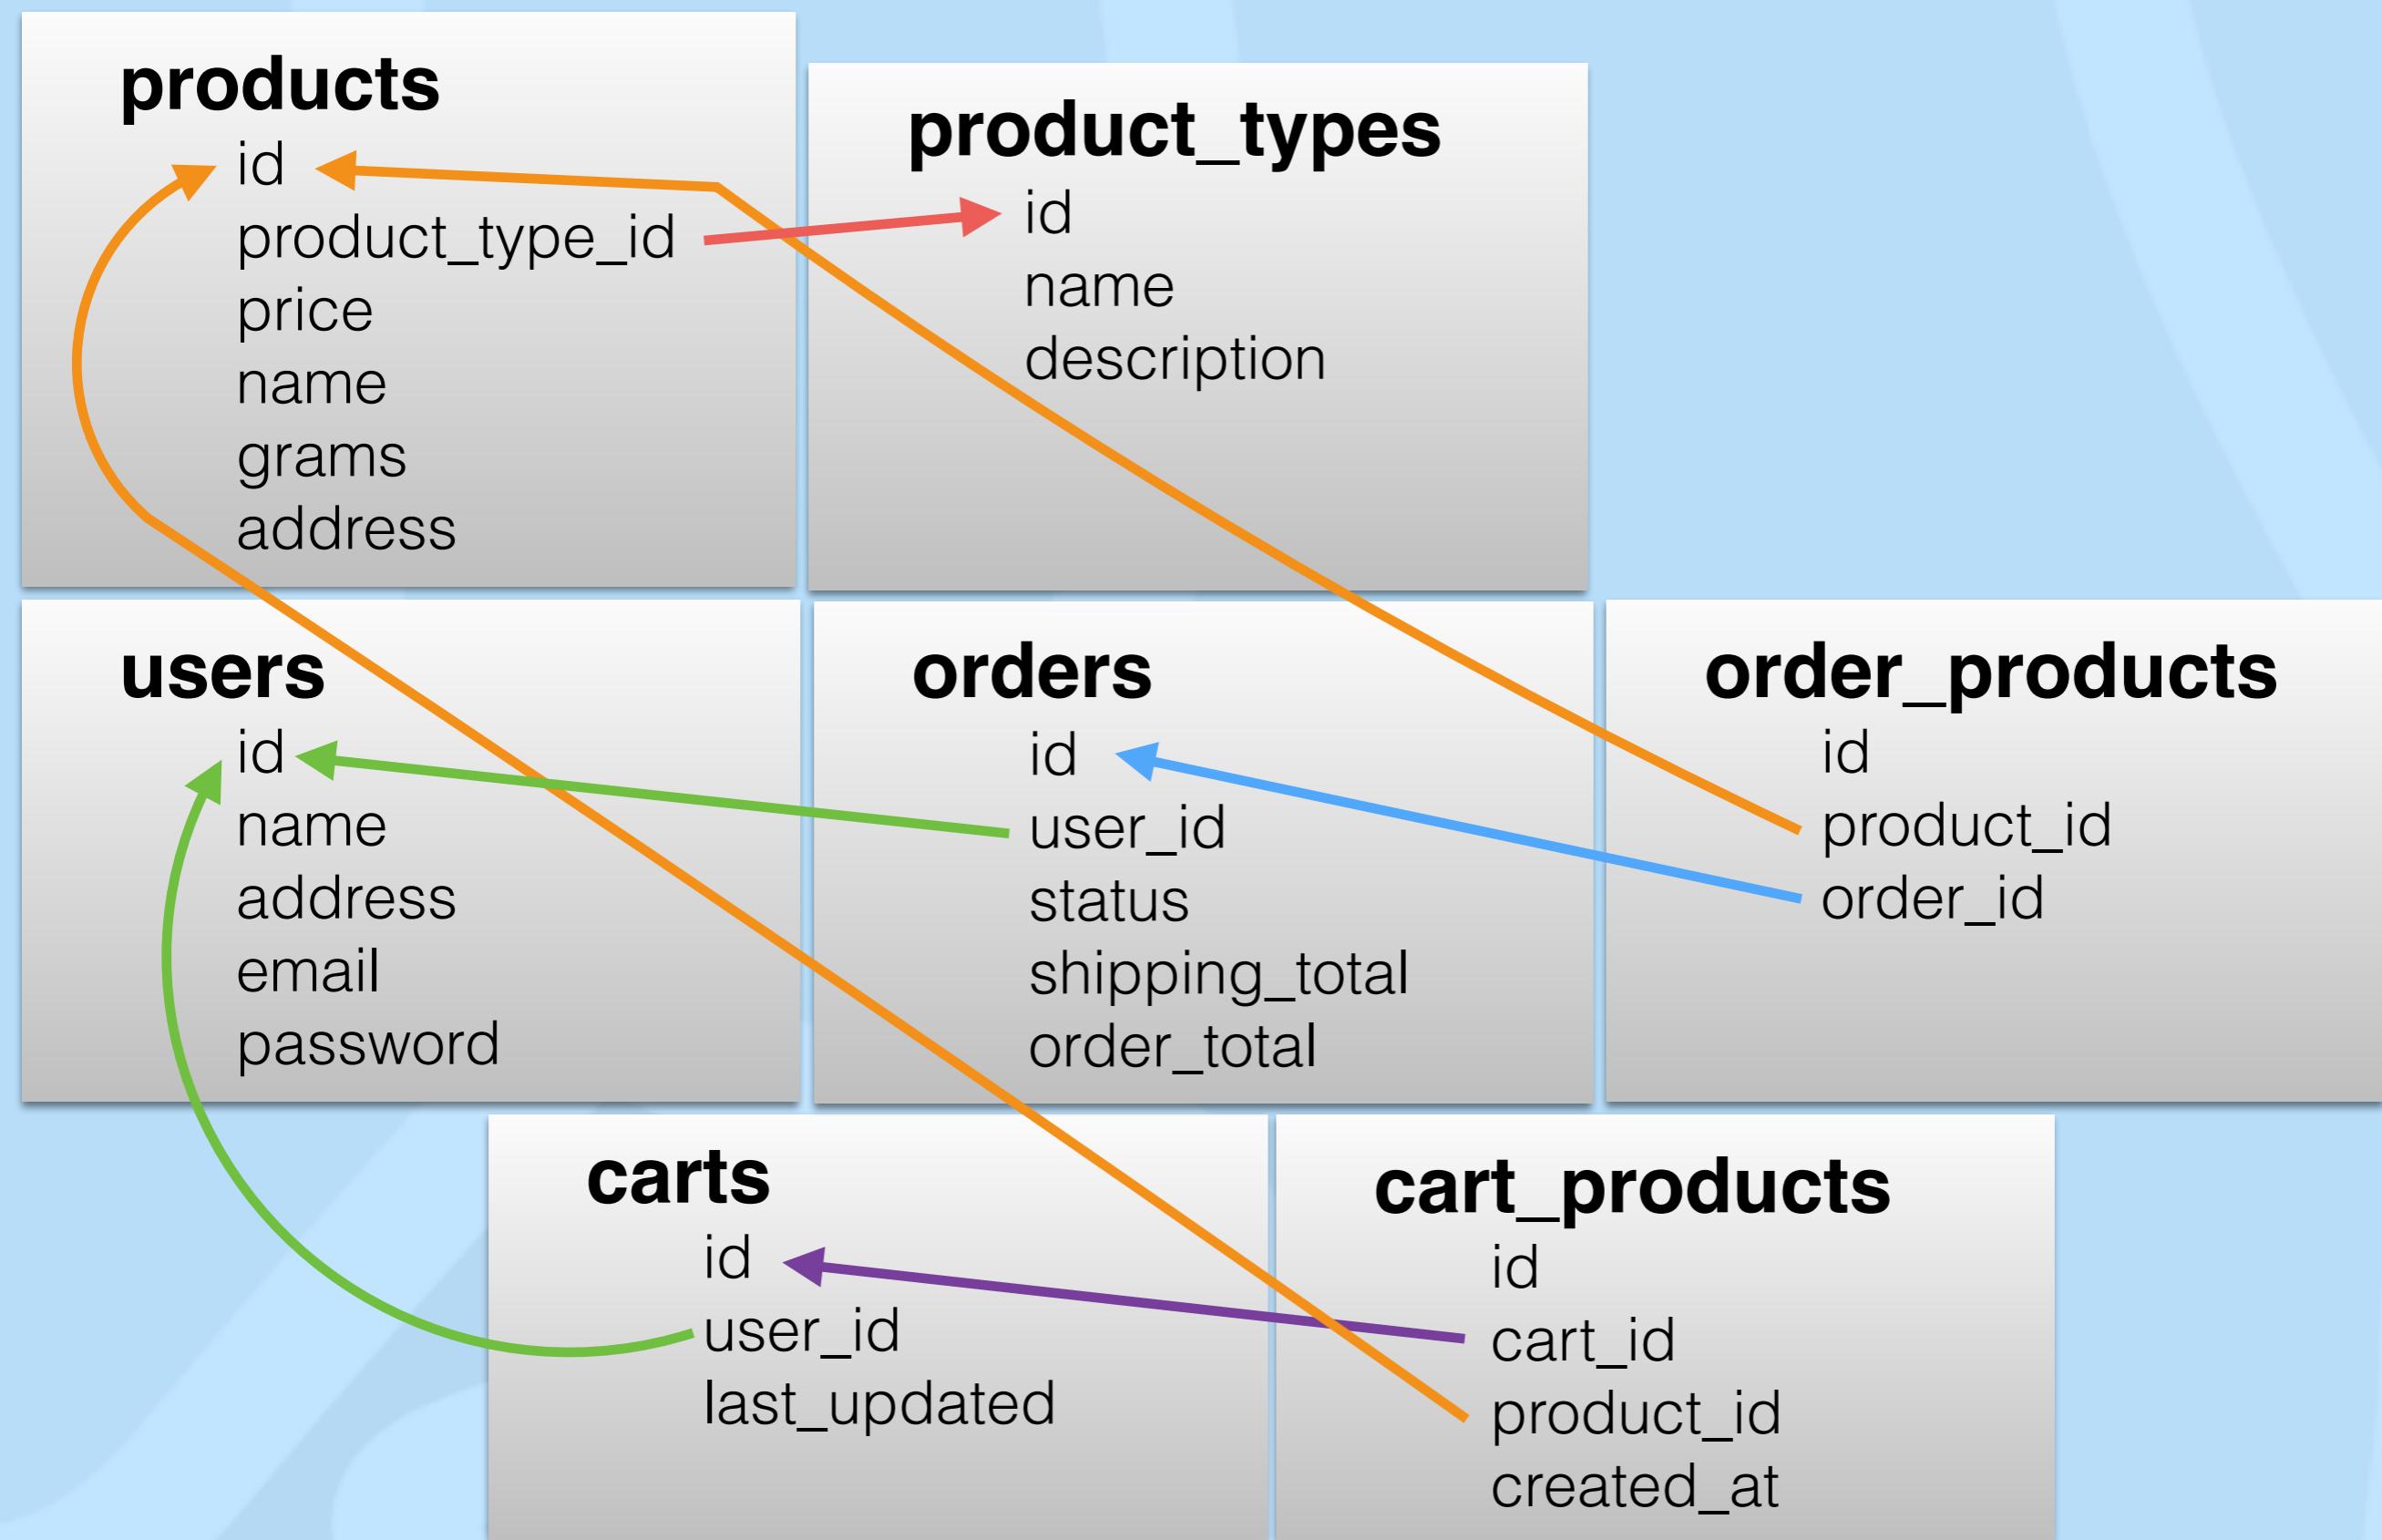
# Case Study: Shopping site



# Case Study: Shopping site



# Case Study: Shopping site



# Key Concept Review

- Select Functions
- Casting & data types
- Sets
- Creating tables

# Key Concept Review

- select functions
  - distinct and group by
  - math and date functions
  - case statements
  - aliasing

# Key Concept Review

- distinct and group by

```
SELECT DISTINCT product_type_id FROM products ORDER  
BY product_type_id;
```

```
SELECT product_type_id FROM products  
GROUP BY product_type_id  
ORDER BY product_type_id;
```

# Key Concept Review

- math functions: avg, count, sum; date functions
  - `SELECT SUM(shipping_total) as shipping_total,  
COUNT(*) as total_orders, AVG(order_total) as  
avg_order,  
EXTRACT(MONTH FROM created_at) as month,  
EXTRACT(YEAR FROM created_at) as year  
FROM orders  
GROUP BY year, month  
ORDER BY year, month;`

# Key Concept Review

- select functions
  - aliasing - setting name for columns, tables and values to be referred to more
    - ex: `SELECT o.shipping_total AS total FROM orders o;`

# Key Concept Review

- select functions
  - aliasing - setting name for columns, tables and values to be referred to more
  - ex: 

```
SELECT o.shipping_total AS total
FROM orders o
JOIN order_products op on op.order_id = o.id
JOIN cart_products cp on cp.product_id =
op.product_id
```

# Casting: Data Types

- Casting values - common 'types'
  - Text - a 'string' of characters, words/characters (i.e. from a keyboard)
  - Integer - whole number
  - Decimal - Float (i.e. floating point)
  - Date - a calendar date
  - Timestamp - a date with specific time, can also have a specific timezone

# Casting

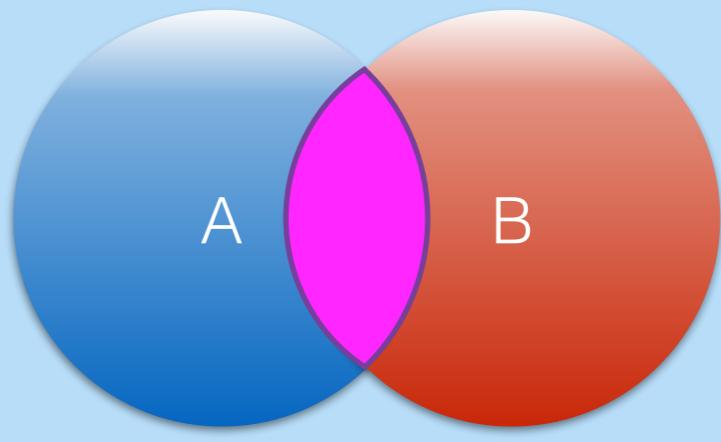
- Casting values
  - 2 methods:
    - **CAST(<column/result> AS <type>)**
    - **<column/result>::<type>**
  - SELECT  
**CAST(o.shipping\_total AS integer)** as int\_total,  
**o.shipping\_total::INTEGER** as int\_total\_v2  
FROM orders o;

# Key Concept Review

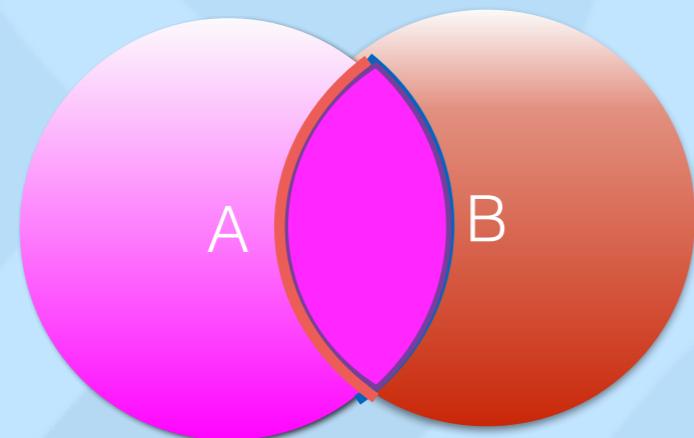
- Sets
  - Outer vs Inner
  - Left vs Right
  - Union, Intersect Except

# Key Concept Review

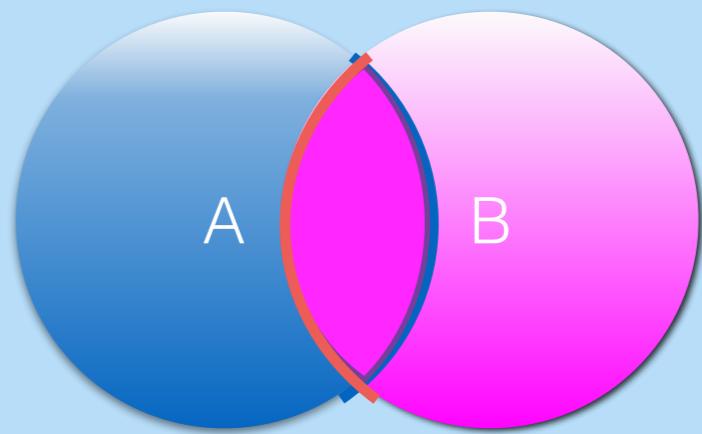
- Sets overview
  - Outer vs Inner
  - Left vs Right



A INNER JOIN B



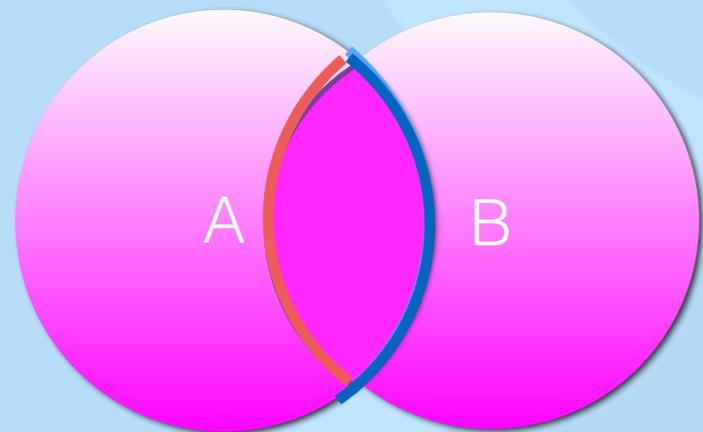
A LEFT OUTER B



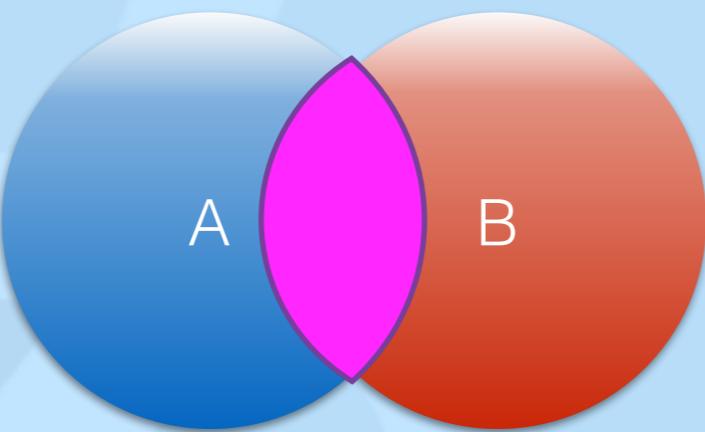
A RIGHT OUTER B

# Key Concept Review

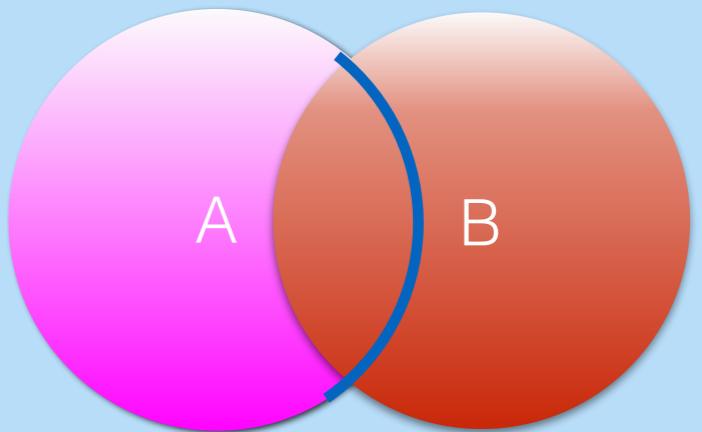
- Sets overview
  - Union, Intersect Except



A union B



A intersect B



A except B

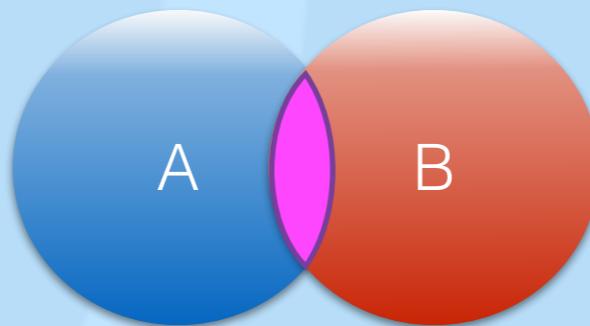
# Inner Join

A = Users

id	integer
name	text
address	text
created_at	date

B = Orders

id	integer
user_id	integer
order_total	float
address	text



```
SELECT * FROM
users INNER JOIN orders ON users.id = orders.user_id
```

AA

id	name	address	created_at	id	user_id	order_total	address
1	John Nob...	123 nowhe...	2016-03-19	8	1	147.33	678 nowhe...
2	John Som...	234 a pla...	2016-05-23	23	2	43.27	372 other...
3	John Doe	345 Doesv...	2016-05-22	52	3	57.9	345 Doesv...

BB1

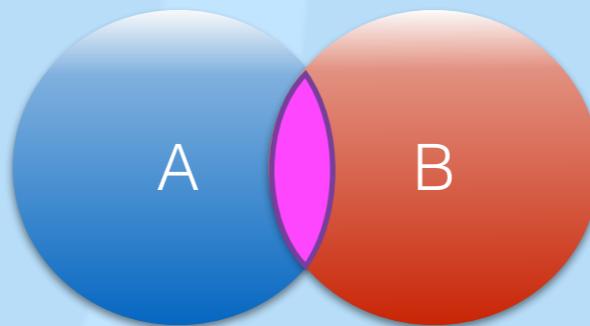
# Inner Join

A = Users

id	integer
name	text
address	text
created_at	date

B = Orders

id	integer
user_id	integer
order_total	float
address	text



```
SELECT * FROM
    users, orders
WHERE users.id = orders.user_id
```

A

	name	address	created_at			user_id	order_total	address
1	John Nob...	123 nowhe...	2016-03-19	8		1	147.33	678 nowhe...
2	John Som...	234 a pla...	2016-05-23	23		2	43.27	372 other...
3	John Doe	345 Doesv...	2016-05-22	52		3	57.9	345 Doesv...

B-1

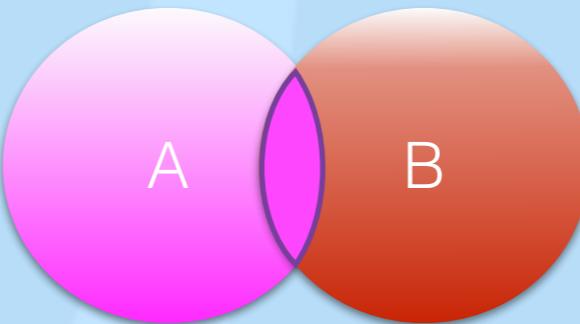
# Left Outer Join

A = Users

id	integer
name	text
address	text
created_at	date

B = Orders

id	integer
user_id	integer
order_total	float
address	text



```
SELECT * FROM users
LEFT OUTER JOIN orders ON users.id = orders.user_id
```

A

id	name	address	created_at		id	user_id	order_total	address
1	John Nob...	123 nowhe...	2016-03-19		8	1	147.33	678 nowhe...
2	John Som...	234 a pla...	2016-05-23		23	2	43.27	372 other...
3	John Doe	345 Doesv...	2016-05-22		52	3	57.9	345 Doesv...
14	Jane Smith	502 downt...	2017-10-21	null	null		null	null
17	Janet Farm	402 lastt...	2018-01-24	null	null		null	null

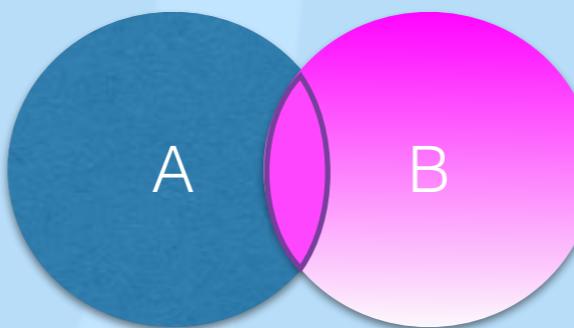
# Right Join

A = Users

id	integer
name	text
address	text
created_at	date

B = Orders

id	integer
user_id	integer
order_total	float
address	text



```
SELECT * FROM users
RIGHT OUTER JOIN orders ON users.id = orders.user_id
```

A

id	name	address	created_at
1	John Nob...	123 nowhe...	2016-03-19
2	John Som...	234 a pla...	2016-05-23
3	John Doe	345 Doesv...	2016-05-22

B

id	user_id	order_total	address
8	1	147.33	678 nowhe...
23	2	43.27	372 other...
52	3	57.9	345 Doesv...

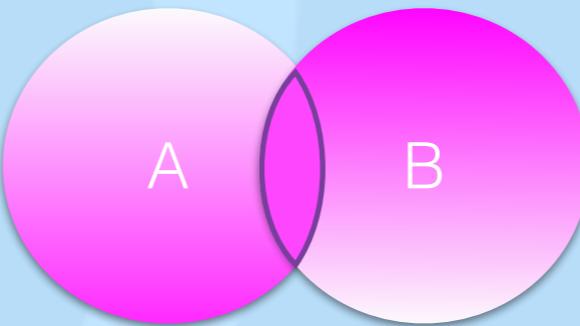
# Full Join

A = Users

<code>id</code>	<code>integer</code>
<code>name</code>	<code>text</code>
<code>address</code>	<code>text</code>
<code>created_at</code>	<code>date</code>

B = Orders

<code>id</code>	<code>integer</code>
<code>user_id</code>	<code>integer</code>
<code>order_total</code>	<code>float</code>
<code>address</code>	<code>text</code>



Users FULL OUTER JOIN orders ON users.address =  
orders.address

A

B

<code>id</code>	<code>name</code>	<code>address</code>	<code>created_at</code>	<code>id</code>	<code>user_id</code>	<code>order_total</code>	<code>address</code>
1	John Nob...	123	2016-03-19	null	null	null	null
14	Jane Smith	502	2017-10-21	null	null	null	null
2	John Som...	234 a	2016-05-23	35	2	113.99	234 a pla...
3	John Doe	345	2016-05-22	52	3	57.9	345 Doesv..
4	John Moe	456	2016-05-22	61	4	74.99	456 somew..
null	null	null	null	109	9	19.24	36127 av..
null	null	null	null <sup>35</sup>	20	1	64.23	678 nowh..

# Advanced Functions

- Having
- String Functions
- Random

# Having

- Having allows you to filter the results on the result of the rest of the query
  - Useful for filtering results of aggregate and math results only
    - Won't work with case statement results for example

# Having

- Having allows you to filter the results on the RESULT of the rest of the query
- WHERE  
Select product\_id,  
count(\*) from  
order\_products  
**WHERE count(\*) > 2**  
GROUP BY product\_id  
ORDER BY count desc
- HAVING  
Select product\_id,  
count(\*) from  
order\_products  
GROUP BY product\_id  
**HAVING count(\*) > 2**  
ORDER BY count(\*) desc

# Advanced functions con't.

- String Functions
  - `||`, concatenate

# Advanced functions con't.

- String Functions
  - `||`, concatenate - push two text values together

# Advanced functions con't.

- String Functions
  - `||`, concatenate - push two text values together

Example: `select 'add text: ' || name from products;`

# Concatenate: Quick Exercise

- String Functions
  - `||`, concatenate - push two text values together

Example: select 'add text' || name from products;

- EXERCISE: Return all users' full address with their name and address concatenated together.

# Concatenate: Quick Exercise

- String Functions
  - `||`, `concatenate` - push two text values together

Example: `select 'add text' || name from products;`

- EXERCISE: Return all users' full address with their name and address concatenated together.

`SELECT name || ' ' || address from users;`

# Advanced functions con't.

- Random

# Advanced functions con't.

- Random - returns a random value between 0 and 1  
EX: Select random()

# Advanced functions con't.

- Random - returns a random value between 0 and 9  
EX: Select random() \* 10

# Advanced functions con't.

- Random - returns a random value between 0 and 1  
EX: Select random() \* 10
- Getting whole numbers
  - round - round to the nearest whole number
  - floor - round down to the nearest whole number less than the value
  - ceiling - round up to the nearest whole number above the value

# Advanced functions con't.

- Random - returns a random value between 0 and 1  
EX: Select random() \* 10
- Getting whole numbers
  - round - round to the nearest whole number
  - floor - round down to the nearest whole number less than the value
  - ceiling - round up to the nearest whole number above the value
- SELECT random\_value,  
round(random\_value),  
floor(random\_value),  
ceiling(random\_value)  
FROM (select random() \* 10 AS random\_value ) random\_query

ADD SPECIFIC NUMBER  
RANDOM W/OUT \* 10

# Random: Quick Exercise

- Random - returns a random value between 0 and 1  
EX: Select random()
- Write a query to return a random whole number between 5 and 10

# Random: Quick Exercise

- Random - returns a random value between 0 and 1  
EX: Select random()
- Write a query to return a random whole number between 5 and 10  
**(HINT: SELECT 5 + <random whole # from 0-5> )**

# Random: Quick Exercise

- Random - returns a random value between 0 and 1  
EX: Select random()
- Write a query to return a random value between 5 and 10

```
SELECT 5 + round(random() * 5)
```

# PgAdmin4 Overview

- PG Admin4 intro

MOVE ME!



**Q & A**

# Exercise: Advanced functions

1. Write a query to return each product name with a random whole number between 1 and 10 concatenated to product name  
(HINT: Use products table;  
|| and random() functions)

# Exercise: Advanced functions

1. Write a query to return each product name with a random number between 1 and 10 concatenated to product name  
(HINT: Use products table, || and RANDOM)

```
select name || ' ' || (1 + round(random() * 9) )  
from products;
```

```
select name || ' ' || (ceiling(random() * 10) )  
from products;
```

# Performance checks & PGAdmin 4

# Performance checks & PGAdmin 4

- Explain

# Performance checks & PGAdmin 4

- Explain
  - Nested
  - Hash
    - Hash full
  - sort
  - aggregate



**Q & A**



# 10 Min Break

# Sub-select, With and temporary tables

- Help to organize your data
- Reduce cognitive load
- Separate process components
- Improve performance

# Sub-select Commands

# Sub-select Commands

- running a query within '(' and ')' and utilizing those results within a query

# Sub-select Commands

- running a query within '(' and ')' and utilizing those results within a query
- EX: Round a single random value:

```
SELECT random_value,  
round(random_value),  
floor(random_value),  
ceiling(random_value)  
FROM (select random() * 10 AS random_value )  
random_query
```

# Sub-select Commands

- running a query within '(' and ')' and utilizing those results within a query
- EX: Get all products with product types that are 'clothing' using a sub-select.

# Sub-select Commands

- running a query within ( and ) and utilizing those results within a query
- EX: Get all products with product types that are 'clothing' using a sub-select.
- Step 1: create a query to get all the product type IDs that are clothing.

# Sub-select Commands

- **Sub-select with IN**

```
SELECT * FROM products  
WHERE product_type_id IN (select id from  
product_types where name ilike '%clothing');
```

# Sub-select Commands

- **Where** with And, **Or**, **In**

return products that are either heavy or products for the home

```
SELECT * FROM products
WHERE grams > 10
OR product_type_id IN (select id from products
where name = 'home decor' OR name = 'kitchen
and dining');
```

# Performance checks & PGAdmin 4

```
1 SELECT * FROM products
2 WHERE grams > 10
3 OR product_type_id IN
4 (select id from products where name = 'home decor' OR name = 'kitchen and dining')
```

Data Output Explain Messages Query History

SubPlan 1

products

products

Filter	((name)::text = 'home decor'::text) OR ((name)::text = 'kitchen and dining'::text)
Node Type	Seq Scan
Relation Name	products
Alias	products_1
Parallel Aware	false
Parent Relationship	SubPlan
Subplan Name	SubPlan 1

# Sub-select Commands

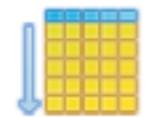
- **Sub-select with JOINS**

```
SELECT * FROM products p
JOIN (select * from product_types where name ilike
'%clothing') p_types ON p_types.id =
p.product_type_id;
```

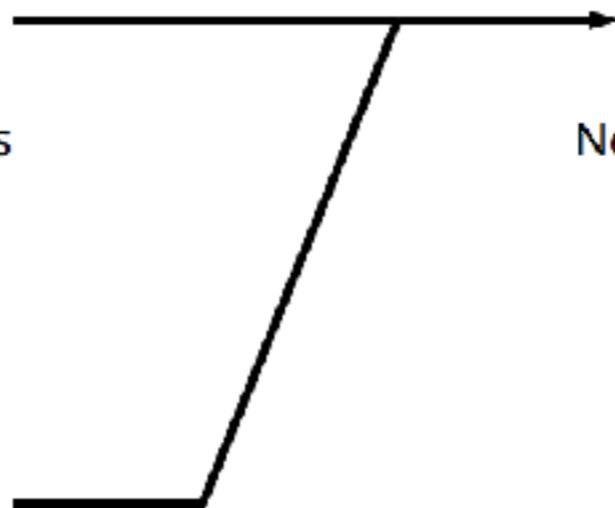
# Performance checks & PGAdmin 4

```
1 SELECT * FROM products p
2 JOIN (select * from product_types
3       where name ilike '%clothing') p_types ON p_types.id = p.product_type_id;
```

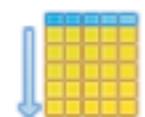
Data Output Explain Messages Query History



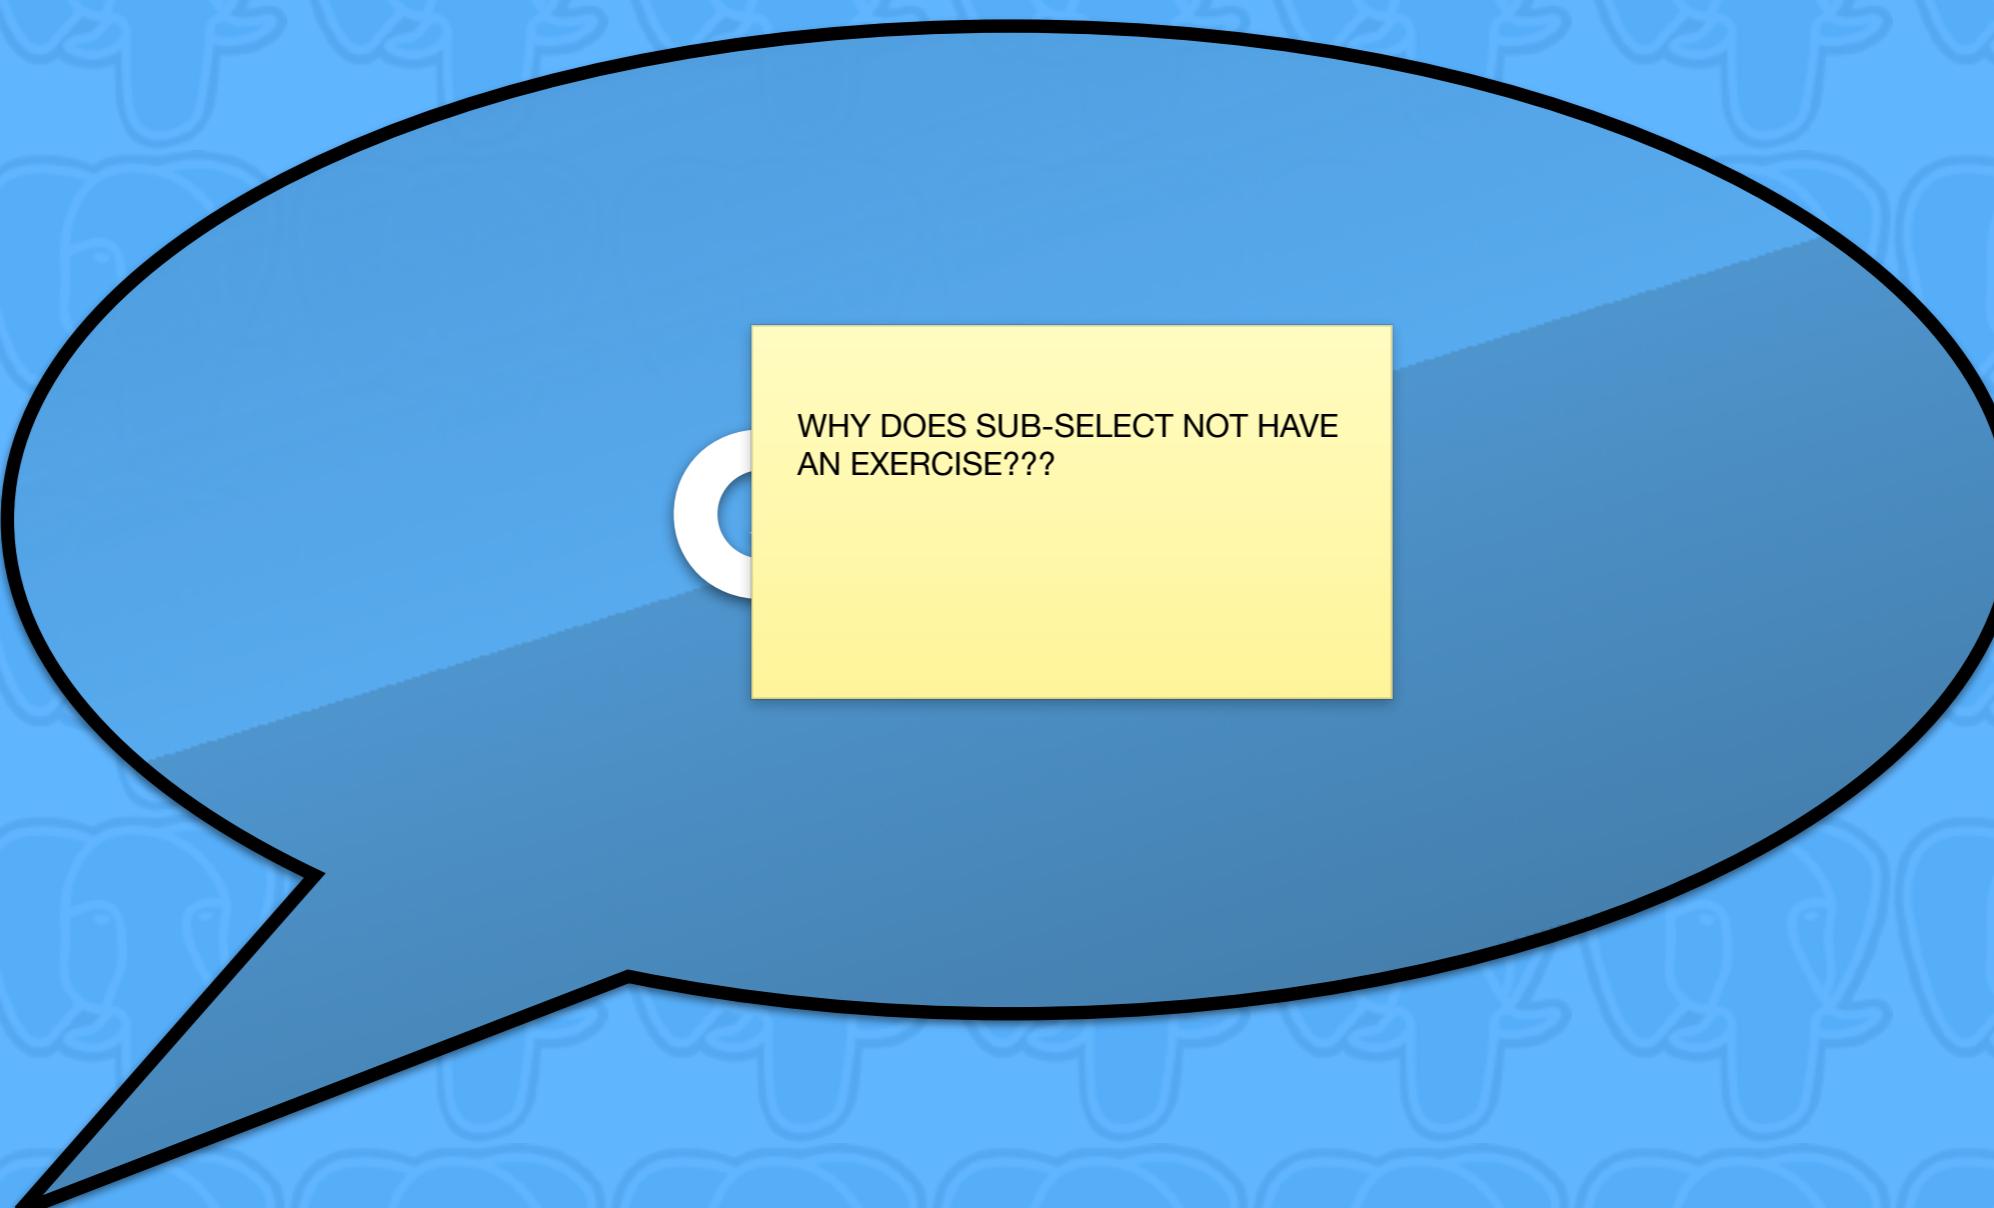
product\_types



Nested Loop Inner  
Join



products



WHY DOES SUB-SELECT NOT HAVE  
AN EXERCISE???

# With

- **WITH (a.k.a. common table expressions (CTE) )**

Temporary view of the data only available in the scope of the query it runs in

- Named so can be interacted with in queries as tables
- fundamentally equivalent to sub-select statements (also CTEs)
- FORMAT:
  - `WITH use_query_as_table as (<query from database>)`
  - `<Query utilizing temp_view>;`

# With

- **WITH**

Temporary view of the data only available in the scope of the query it runs in

- Named so can be interacted with in queries as tables

- FORMAT:

- `WITH use_query_as_table as (<query from database>)`

```
<Query utilizing temp_view>;
```

# With

- **WITH**

Temporary view of the data only available in the scope of the query it runs in

- Named so can be interacted with in queries as tables

- FORMAT:

- `WITH use_query_as_table as (<query from database>)  
<Query utilizing temp_view>;`

- **Example:**

**WITH clothing\_product\_types as (select id from product\_types  
where name ilike '%clothing')**

**SELECT \* FROM products**

**WHERE product\_type\_id IN (select id from clothing\_product\_types);**

# With: Quick Exercise

- Use WITH to create a set of user\_ids that have updated their cart this year;  
select all user information for user who have updated their cart this year  
(HINT: with will select from carts, the select statement will select from users)

format:

```
WITH use_query_as_table as (<query from database>)
<Query utilizing temp_view>;
```

# With: Quick Exercise

- Use WITH to create a set of user\_ids that have updated their cart this year; select all user information for user who have updated their cart this year  
(HINT: with will select from carts  
last\_updated > '2018-01-01',  
the select statement will select from users)

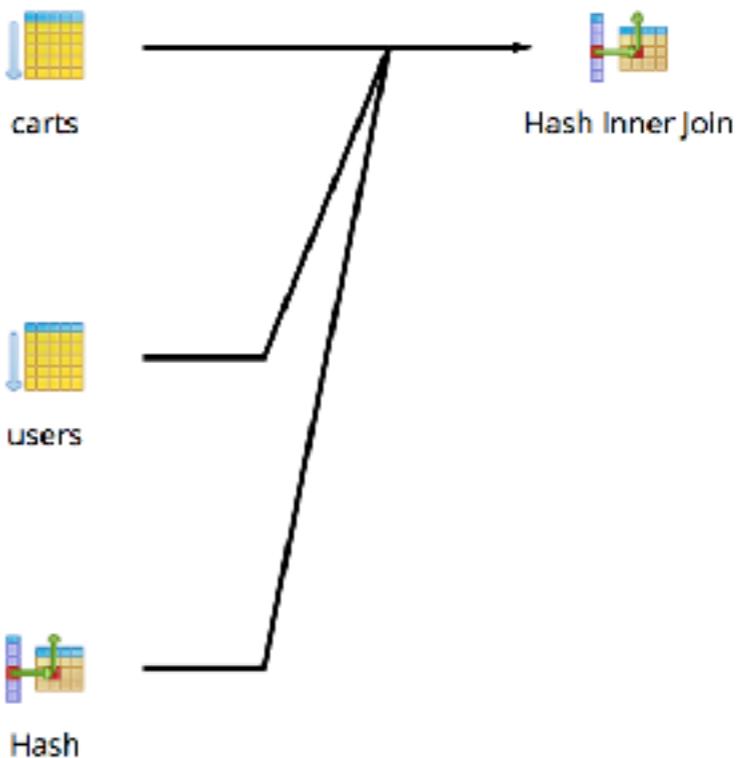
```
WITH user_updated_carts AS (SELECT user_id from
carts where last_updated > '2018-01-01')
```

```
SELECT * FROM users
WHERE id IN (select user_id from
user_updated_carts)
```

# Performance checks & PGAdmin 4

```
1 WITH user_updated_carts AS (SELECT user_Id from carts  
2 where last_updated > '2018-01-01')  
3  
4 SELECT * FROM users  
5 WHERE id IN (select user_id from  
6 user_updated_carts)
```

Data Output Explain Messages Query History



# With performance

- **With** and **IN**

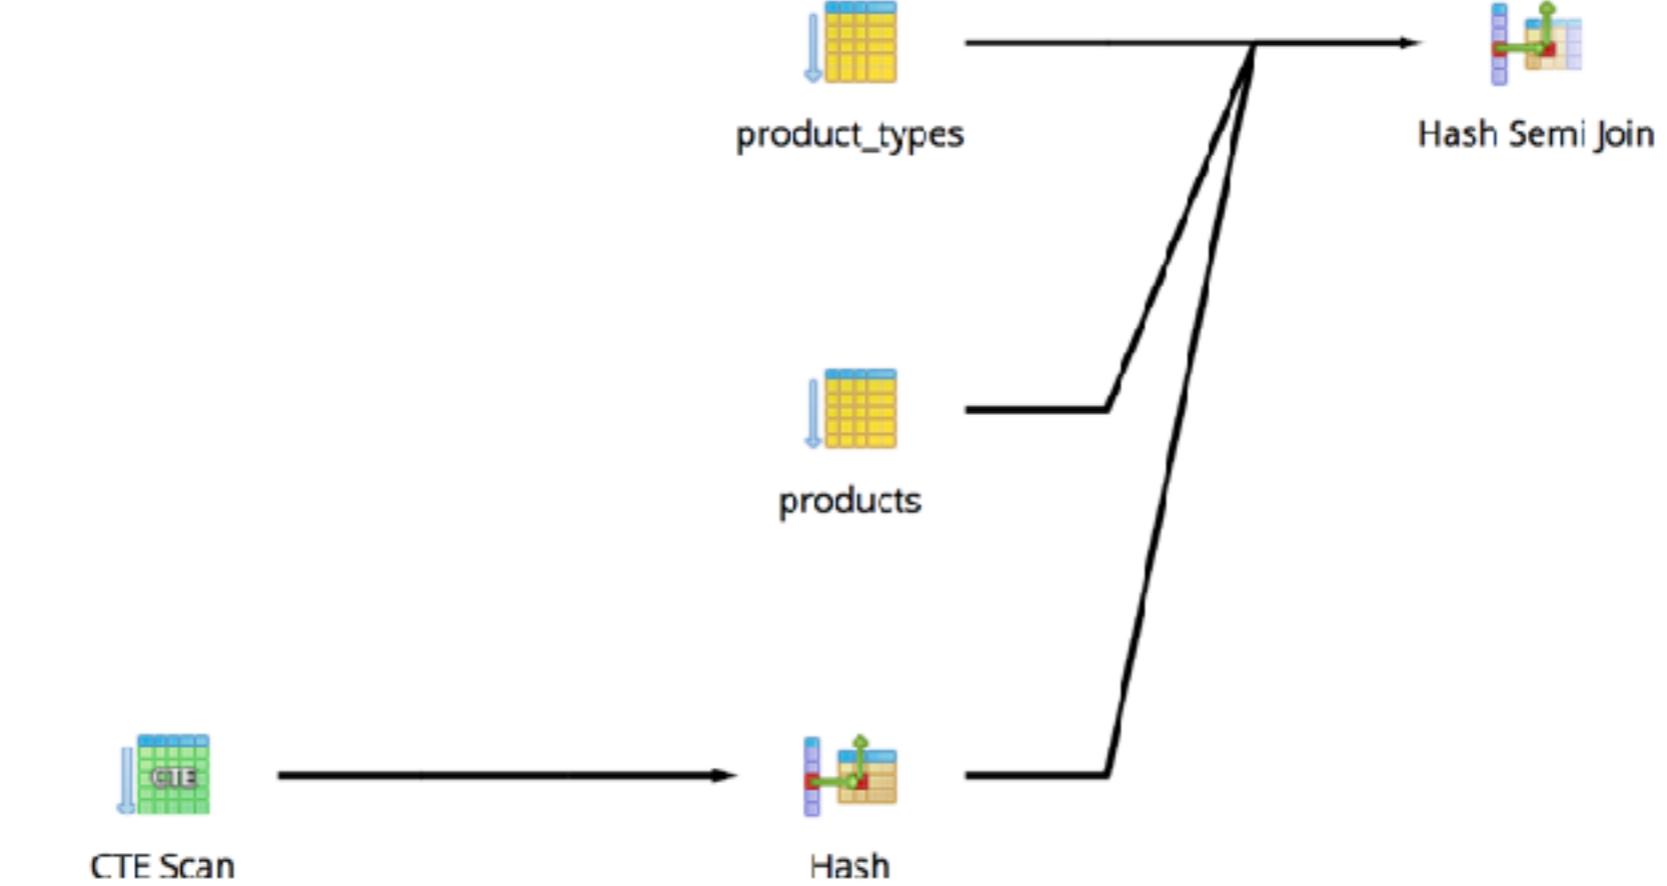
```
WITH clothing_product_types as (select id from
product_types where name ilike '%clothing')
```

```
SELECT * FROM products
WHERE product_type_id IN (select id from
clothing_product_types);
```

# Performance checks & PGAdmin 4

```
1 WITH clothing_product_types AS (SELECT id FROM product_types WHERE name ilike '%clothing')
2
3 SELECT * FROM products
4 WHERE product_type_id IN (SELECT id FROM clothing_product_types);
```

Data Output Explain Messages Query History



# With

- **WITH** and **JOIN**

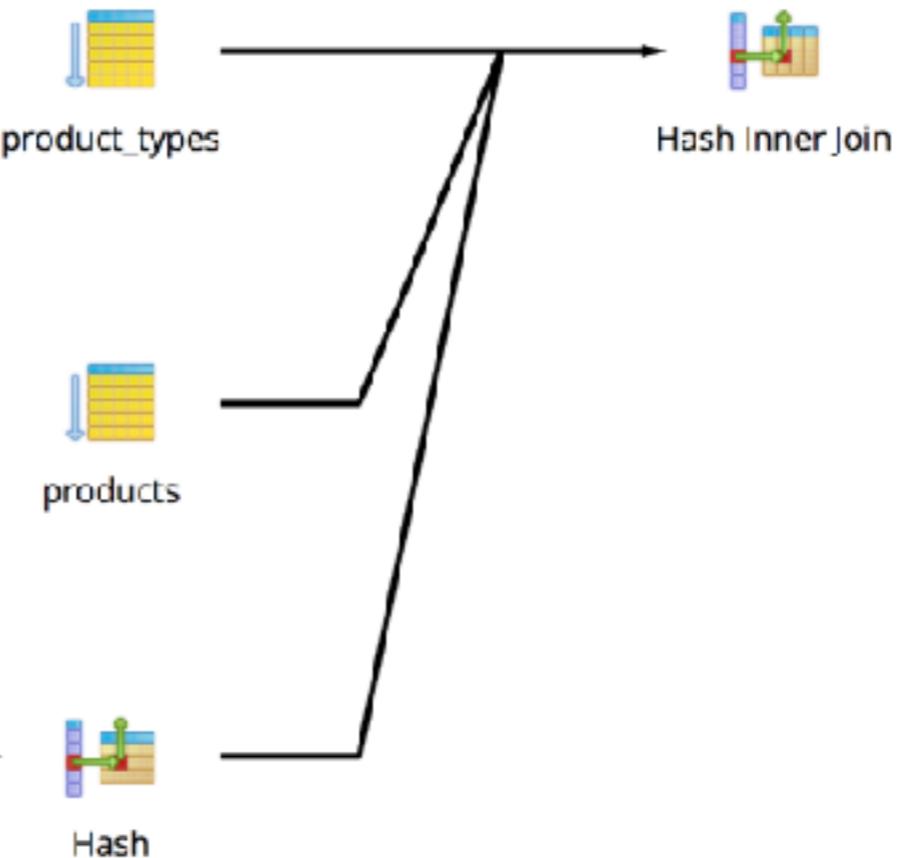
```
WITH clothing_product_types as (select * from  
product_types where name ilike '%clothing')
```

```
SELECT * FROM products  
JOIN clothing_product_types p_types ON  
p_types.id = p.product_type_id;
```

# Performance checks & PGAdmin 4

```
1 WITH clothing_product_types AS (SELECT * FROM product_types WHERE name ilike '%clothing')
2
3 SELECT * FROM products
4 JOIN clothing_product_types p_types ON p_types.id = products.product_type_id;
```

Data Output Explain Messages Query History





**Q & A**

# Temporary Tables

# Temporary Tables

- Temporary Tables
  - provide similar simplification to complex queries,
  - Performance optimization control
  - Format: CREATE TEMPORARY TABLE <table\_name> as (<query>); SELECT \* FROM <table\_name>;

# Temporary Tables

- Based on:  
WITH clothing\_product\_types as (select \* from product\_types where name ilike '%clothing')  
  
SELECT \* FROM products  
JOIN clothing\_product\_types p\_types ON p\_types.id = p.product\_type\_id;

FORMAT - CREATE TEMPORARY TABLE <table\_name> as (<query>); SELECT \* FROM <table\_name>;

# Temporary Tables

- Based on:

```
WITH clothing_product_types as (select * from product_types where name ilike '%clothing')
```

```
SELECT * FROM products  
JOIN clothing_product_types p_types ON p_types.id = p.product_type_id;
```

**FORMAT** - CREATE TEMPORARY TABLE <table\_name> as (<query>); SELECT \* FROM <table\_name>;

```
CREATE TEMPORARY TABLE clothing_product_types as (select *  
from product_types where name ilike '%clothing');
```

```
SELECT * FROM products  
JOIN clothing_product_types p_types ON p_types.id =  
product_type_id;
```

# Temporary Tables: Quick Exercise

- Based on the with statement from earlier. Use a TEMPORARY TABLE to hold all users/user\_ids that have updated their cart this year. Select all user information from the users table for those users:

Based ON: WITH user\_updated\_carts AS (SELECT user\_Id from carts where last\_updated > '2018-01-01')  
SELECT \* FROM users where user\_id in (user\_updated\_carts);

**FORMAT** - CREATE TEMPORARY TABLE <table\_name> as (<query>); SELECT \* FROM <table\_name>;

# Temporary Tables: Quick Exercise

- Based on the with statement from earlier. Use a TEMPORARY TABLE to hold all users/user\_ids that have updated their cart this year. Select all user information from the users table for those users:  
Based ON: WITH user\_updated\_carts AS (SELECT user\_id from carts where last\_updated > '2018-01-01')

```
CREATE TEMPORARY TABLE user_updated_carts as (select * from carts where last_updated >= '2018-01-01');
```

```
SELECT * FROM users  
WHERE id IN (select user_id from  
user_updated_carts);
```

# Performance checks & PGAdmin 4

```
1 CREATE TEMPORARY TABLE clothing_product_types AS (SELECT * FROM product_types  
2 WHERE name ilike '%clothing')  
3 ;  
4 EXPLAIN  
5 SELECT * FROM products  
6 JOIN clothing_product_types p_types ON p_types.id = products.product_type_id;
```

Data Output Explain Messages Query History

## QUERY PLAN

text

1 Hash Join (cost=2.24..19.92 rows=88 width=257)

2 Hash Cond: (p\_types.id = products.product\_type\_id)

3 -> Seq Scan on clothing\_product\_types p\_types (cost=0.00..13.20 rows=320 width=222)

4 -> Hash (cost=1.55..1.55 rows=55 width=35)

5 -> Seq Scan on products (cost=0.00..1.55 rows=55 width=35)



**Q & A**

# With, Sub-select, and Temporary Tables Exercises

- Write a query to randomly assign each user a 'bucket' that is a value between 1 and 3. Return the total number of orders for users in each bucket.

Use each of the following to store this result as user\_buckets

***(SELECT id as user\_id, 1 + round(random() \* 2) AS bucket  
FROM users)***

- 1) Sub-select - Using '(' and ')'
- 2) WITH
- 3) TEMPORARY TABLE

# With, Sub-select, and Temporary Tables Exercises

- Write a query to randomly assign each user a 'bucket' that is a value between 1 and 3. Return the total number of orders for users in each bucket.

Use each of the following to store this result as user\_buckets:

**(*SELECT id as user\_id, 1 + round(random() \* 2) AS bucket  
FROM users*)**

**HINT:** Join the results and group by user\_ID:

*JOIN orders ON orders.user\_id = user\_buckets.user\_id*

- 1) Sub-select - Using '(' and ')'
- 2) WITH
- 3) TEMPORARY TABLE

# With, Sub-select, and Temporary Tables Exercises

- Write a query to randomly assign each user a 'bucket' that is a value between 1 and 3. Return the total number of orders for users in each bucket.

Use each of the following to store this result as user\_buckets:

**(SELECT id as user\_id, 1 + round(random() \* 2) AS bucket FROM users)**

**HINT:** Join the results and group by user\_ID:

SELECT bucket, count(\*) as total\_orders...

...

*JOIN orders ON orders.user\_id = user\_buckets.user\_id  
GROUP BY bucket*

- 1) Sub-select - Using '(' and ')'
- 2) WITH
- 3) TEMPORARY TABLE

# With, Sub-select, and Temporary Tables Exercises

- Write a query to randomly assign each user a 'bucket' that is a value between 1 and 3. Return the total number of orders for users in each bucket.

Use each of the following to store this result as user\_buckets

**(SELECT id, 1 + round(random() \* 2) AS bucket FROM users)**

**HINT:** Join the results and group by user\_ID:

*JOIN orders ON orders.user\_id = user\_buckets.user\_id*

1) Sub-select - Using '(' and ')'

*SELECT bucket, count(orders.id) FROM orders*

*JOIN (Select id, 1 + round(random() \* 2) as bucket from users) user\_buckets on user\_buckets.id = orders.user\_id*

*GROUP BY BUCKET*

# With, Sub-select, and Temporary Tables Exercises

- Write a query to randomly assign each user a 'bucket' that is a value between 1 and 3. Return the total number of orders for users in each bucket.

Use each of the following to store this result as user\_buckets

**(SELECT id, 1 + round(random() \* 2) AS bucket FROM users)**

**HINT:** Join the results and group by user\_ID:

*JOIN orders ON orders.user\_id = user\_buckets.user\_id*

2) WITH

**WITH user\_buckets AS (Select id, 1 + round(random() \* 2) as bucket from users)**

*SELECT bucket, count(orders.id) FROM orders*

*JOIN user\_buckets on user\_buckets.id = orders.user\_id*

*GROUP BY BUCKET*

# With, Sub-select, and Temporary Tables Exercises

- Write a query to randomly assign each user a 'bucket' that is a value between 1 and 3. Return the total number of orders for users in each bucket.

Use each of the following to store this result as user\_buckets

**(SELECT id, 1 + round(random() \* 2) AS bucket FROM users)**

**HINT:** Join the results and group by user\_ID:

*JOIN orders ON orders.user\_id = user\_buckets.user\_id*

3) Temporary Table

**CREATE TEMPORARY TABLE** user\_buckets **AS** (**Select id, 1 + round(random() \* 2) as bucket from users;**)

**SELECT bucket, count(orders.id) FROM orders**

*JOIN user\_buckets on user\_buckets.id = orders.user\_id*

**GROUP BY BUCKET**



**Q & A**



# 10 Min Break

# Views

# Views

- Persisted outside of a single session or query

# Views

- Persisted outside of a single session or query
- Reflect updates to tables used by views - query ran each time view is referenced

# Views

- Persisted outside of a single session or query
- Reflect updates to tables used by views - query ran each time view is referenced
- Materialized Views - Store the results on disk

# Views

- Persisted outside of a single session or query
- Reflect updates to tables used by views - query ran each time view is referenced
- Materialized Views - Store the results on disk
- Can also be temporary

# Views

- Common Uses
  - Summarize/consolidate data
  - Re-structure data to be easier to consume/use
  - Utilize separate permissions to limit the data some users have access to

# Creating Views

- Create

# Creating Views

- Create  
CREATE VIEW <view\_name> AS <query>

# Creating Views

- Create  
CREATE VIEW <view\_name>  
(<col1\_name>, <col2\_name>, ...)  
AS <query>

# Creating Views

- `CREATE VIEW <view_name> AS <query>`
  - `CREATE VIEW <view_name> (<col1_name>, <col2_name>, ...) AS <query>`
- Ex: `CREATE VIEW clothing AS SELECT id as product_type_id, name FROM product_types WHERE name ilike '%clothing'`
- Ex: `CREATE VIEW clothing (product_type_id, name) AS SELECT id, name FROM product_types WHERE name ilike '%clothing'`

# Creating Views

- Create
- Create or Replace
  - Replace any existing views with the same name if they exist

# Views

- CREATE VIEW <view\_name>  
(<col1\_name>, <col2\_name>, ...) -- *optional*  
AS <query>
- Create or Replace View  
**CREATE OR REPLACE** VIEW <view\_name>  
AS <query>

# Creating Views

- **CREATE OR REPLACE VIEW** <view\_name>  
(<col1\_name>, <col2\_name>, ...) -- *optional*  
AS <query>

Ex: CREATE OR REPLACE VIEW clothing  
AS SELECT id as product\_type\_id,  
name  
FROM product\_types  
WHERE name ilike '%clothing'

# Creating Views

- Create
- Create or Replace
  - Replace any existing views with the same name if they exist
- Temporary Views
  - View only exists during the session

# Creating Views

- CREATE OR REPLACE **TEMPORARY** VIEW  
<view\_name>  
(<col1\_name>, <col2\_name>, ...) -- *optional*  
AS <query>

Ex: CREATE OR REPLACE TEMPORARY VIEW  
temp\_clothing  
AS SELECT id as product\_type\_id,  
name  
FROM product\_types  
WHERE name ilike '%clothing'

# Creating Views

- CREATE OR REPLACE **TEMPORARY** VIEW

```
temp_clothing  
AS SELECT id as product_type_id, name  
FROM product_types  
WHERE name ilike '%clothing'
```

Ex:

```
SELECT * FROM temp_clothing
```

# Creating Views

- CREATE **TEMPORARY** VIEW temp\_tech  
AS SELECT id as product\_type\_id, name  
FROM product\_types  
WHERE name ilike '**computer%**'

Ex:

```
SELECT * FROM temp_tech
```

# Creating Views: Quick Exercise

- Use a VIEW to hold all users/user\_ids that have updated their cart this year. Select all user information from the users table for those users:

Based ON: WITH user\_updated\_carts AS (SELECT user\_Id from carts where last\_updated > '2018-01-01')

CREATE VIEW <view\_name>  
(<col1\_name>, <col2\_name>, ...) -- *optional*  
AS <query> ;

SELECT \* FROM <view\_name>;

# Creating Views: Quick Exercise

- Use a VIEW to hold all users/user\_ids that have updated their cart this year. Select all user information from the users table for those users:

Based ON: WITH user\_updated\_carts AS (SELECT user\_Id from carts where last\_updated > '2018-01-01')

```
CREATE VIEW user_updated_carts  
AS select * from carts where last_updated >= '2018-01-01';
```

```
SELECT * from users where id in (SELECT user_id from  
user_updated_carts)
```



**Q & A**

# Views Lab

# Views Lab

- Group Exercise: Create View

# Views Lab

- Group Exercise: Let's create a view to store order\_Id and all product information

# Views Lab

- Group Exercise: Let's create a view to store order\_id, created at date and product\_ID with all the product's information

# Views Lab

- Group Exercise: Let's create a view to store order\_id, order created at date, user who made the order with all the product information included.

```
CREATE VIEW order_product_details AS
SELECT order_id, product_id,
name as product_name,
price, grams, product_type_id,
orders.created_at AS order_date,
user_id
FROM products
JOIN order_products ON order_products.product_id = products.id
JOIN orders on orders.id = order_id;
```

# Views Lab

- Group Exercise: Let's create a view to store order\_id, order created at date, user who made the order with all the product information included.

```
SELECT count(*) FROM order_product_details;
```

```
SELECT * FROM order_product_details ORDER BY  
order_id, product_id;
```

# Views Lab

- Now let's add some data to one of the tables the view uses to see what happens!

# Views Lab

- Group Exercise: Insert Into Source Tables used in View

```
INSERT INTO orders (user_id, status, shipping_total,  
order_total, address, created_at)  
VALUES  
(13, 'order confirmed', (1 + round(random() *1000)),  
10.00, '123 nowhere 12345', now() ),  
(14, 'order confirmed', (1 + round(random() *1000)),  
20.00, '234 somewhere 23456', now() ),  
(15, 'order confirmed', (1 + round(random() *1000)),  
30.00, '345 anywhere 34567', now() )
```

# Views Lab

- Group Exercise: Insert Into second Source Table so our query returns results!

```
SELECT * FROM orders order by created_at desc;  
select count(*) from products  
INSERT INTO order_products (order_id, product_id)  
VALUES  
(160,1),  
(160,2),  
(160,3),  
(161,4),  
(161,5),  
(161,6),  
(162,7),  
(162,8),  
(162,9),  
(162, (1 + round(random() *50)))
```

# Views Lab

- Group Exercise: Select from view again and your results are up to date!

# Views Lab

- Group Exercise: Select from view again and your results are up to date!

```
SELECT * FROM order_product_details  
ORDER BY order_id desc, product_id;
```

# Views Lab

- Group Exercise: Try to drop a table used in the view

DROP TABLE orders;

# Views Lab

- Group Exercise: Try to drop a table used in the view...

**Cant!**

DROP TABLE carts;

ERROR: cannot drop table carts because other objects depend on it

DETAIL: constraint cart\_products\_cart\_id\_fkey on table cart\_products depends on table carts

**view user\_updated\_carts depends on table carts**

HINT: Use DROP ... CASCADE to drop the dependent objects too.

SQL state: 2BP01

# Views Lab: Cascade

- apply to any and all related or dependent objects (tables, views, etc)

# Views Lab: Cascade

- apply to any and all related or dependent objects (tables, views, etc)
- Use case: Have views and need to drop all artifacts  
-> requires view to be re-created

# Views Lab: Cascade

- Group Exercise: First let's create a table and view to work with

# Views Lab: Cascade

- Group Exercise: First let's create a table and view to work with

```
CREATE TABLE carts_cascade AS SELECT *  
FROM carts;
```

# Views Lab: Cascade

- Group Exercise: First let's create a table and view to work with

```
CREATE VIEW updated_carts  
AS select * from carts_cascade WHERE  
last_updated > '2018-01-01';
```

```
SELECT * FROM updated_carts;
```

# Views Lab: Cascade

- Try to drop the parent table.. should see same error

```
DROP TABLE carts_cascade;
```

# Views Lab: Cascade

- Try to drop the parent table.. should see same error

```
DROP TABLE carts_cascade;
```

ERROR: cannot drop table carts\_cascade because other objects depend on it

**DETAIL: view updated\_carts depends on table carts\_cascade**

HINT: Use DROP ... CASCADE to drop the dependent objects too.

SQL state: 2BP01

# Views Lab: Cascade

- Try to drop the parent table.. should see same error

```
DROP TABLE carts_cascade CASCADE;
```

ERROR: cannot drop table carts\_cascade because other objects depend on it

DETAIL: view updated\_carts depends on table carts\_cascade

**HINT: Use DROP ... CASCADE to drop the dependent objects too.**

SQL state: 2BP01

# Views Lab: Cascade

- Group Exercise:  
DROP TABLE carts\_cascade CASCADE;

# Views Lab: Cascade

- Group Exercise:  
DROP TABLE carts\_cascade CASCADE;

*NOTICE: drop cascades to view updated\_carts*

*DROP TABLE*

*Query returned successfully in 73 msec.*

# Views Lab: Cascade

- Group Exercise: Re-setup the table and view

# Views Lab: Cascade

- Group Exercise: Re-setup the table and view  
CREATE TABLE carts\_cascade AS SELECT \*  
FROM carts;

```
CREATE VIEW updated_carts  
AS select * from carts_cascade WHERE  
last_updated > '2018-01-01';
```

```
SELECT * FROM updated_carts;
```

# Views Lab

- Work-around: Re-load tables using TRUNCATE instead

# Views Lab: Truncate

- Truncate delete all rows in a table  
(equivalent to: DELETE FROM <table>)

```
TRUNCATE <table>;
```

# Views Lab: Truncate

- Truncate delete all rows in a table  
(equivalent to: DELETE FROM <table>)

GROUP EXERCISE: TRUNCATE carts\_cascade;

# Views Lab: Truncate

- Truncate delete all rows in a table  
(equivalent to: DELETE FROM <table>)

GROUP EXERCISE: TRUNCATE carts\_cascade;

SELECT \* FROM updated\_carts;

# Views Lab: Truncate

- USE CASES
  - Replace table
  - Views
  - Permissions



**Q & A**



# 10 Min Break

# Indexes

# Indexes

- What is an index (how do they work at high level)

# Indexes

- What is an index (how do they work at high level)
  - create quick-lookup tables Postgres uses (overhead)
  - Essentially a pointer to data in a table

# Indexes

- what is an index (how do they work at high level)
- Over-indexing

# Indexes

- Speed up: Select, Joins and Where
- Slows Down: Update and Insert

# Indexes

- Primary keys - Automatically constrained to be unique with index

# Performance checks & PGAdmin 4

- Query from Views Lab:

```
SELECT order_id, product_id,  
name as product_name,  
price, grams, product_type_id,  
orders.created_at AS order_date,  
user_id  
FROM products  
JOIN order_products ON  
order_products.product_id = products.id  
JOIN orders on orders.id = order_id;
```

# Performance checks & PGAdmin 4

- Query from Views Lab:

First let's add some random data to the tables:

```
INSERT INTO order_products (order_id, product_id)
```

- ```
SELECT floor(random_to((select max(id) from orders )),floor(random_to((select max(id) from products)) )  
FROM generate_series(1,500000);
```

```
INSERT INTO cart_products (cart_id, product_id, created_at)  
SELECT floor(random_to((select max(id) from carts )),floor(random_to((select max(id) from carts)) ), now()  
FROM generate_series(1,500000);
```

# Indexes

- Add index

# Indexes

- Add index

```
CREATE INDEX <index_name> ON <table>  
(<table_column>)
```

# Indexes

- Add index
- Vacuum (analyze) is required for the database to setup the pointers/look-up data.

# Indexes

- Where to apply indexes

# Indexes

- Where to apply indexes
  - nested joins
  - full table scans (on large tables)

# Indexes

- Where to apply indexes: what columns?
  - columns being joined on
  - column values used in where clauses

# Indexes

- Index for Join clause:

```
CREATE INDEX id_products_index ON products  
(id);
```

```
CREATE INDEX product_id_cart_prod_idx ON  
cart_products (product_id);
```

# Indexes

- Index for Where clause:

```
CREATE INDEX name_products_index ON  
products (name)
```

# Review + Q&A

# Review + Q&A

- TODO: SETUP LARGE DB, STUDENTS AREN'T DOING , JUST GET IT READY AND GIVE IT TO THEM!